

# 1. End-to-End Agent Workflow

flowchart LR

A[Ingestion Agent] --> B[Profiling Agent]  
B --> C[EDA & Visualization Agent]  
C --> D[Cleaning Agent]  
D --> E[Feature Engineering Agent]  
E --> F[Model Training Agent]  
F --> G[HPO Agent]  
G --> H[Evaluation Agent]  
H --> I[Explainability Agent]  
I --> J[Reporting & Dashboard Agent]  
J --> K[Deployment & Monitoring Agent]  
K --> L[Feedback & Iteration Agent]  
L --> M[Knowledge Agent]

style A fill:#f9f,stroke:#333,stroke-width:1px  
style C fill:#ff9,stroke:#333,stroke-width:1px  
style J fill:#9ff,stroke:#333,stroke-width:1px

# 2. Agent & Task Breakdown

Agent	Key Tasks	Tools & Services
1. Data Ingestion	<ul style="list-style-type: none"><li>• Connect to sources (CSV/SQL/Delta/ADLS)</li><li>• Validate schema</li><li>• Snapshot raw data</li></ul>	pandas, Azure Data Factory, Blob Storage
2. Data Profiling	<ul style="list-style-type: none"><li>• Compute column types, cardinalities</li><li>• Detect high-null or constant columns</li><li>• Sample value distributions</li></ul>	pandas-profiling, Azure ML DataPrep
3. EDA & Visualization	<ul style="list-style-type: none"><li>• Chart Recommendation (heatmaps, bar, line, scatter, box)</li><li>• Generate static plots (matplotlib)</li><li>• Produce interactive dashboards (Plotly/Dash or Power BI embeds)</li><li>• Time-series decomposition plots</li></ul>	matplotlib, Seaborn via code interpreter  Plotly, Power BI, Streamlit

	<ul style="list-style-type: none"> <li>• Automatic commentary: “Sales peaked in Q4...” via OpenAI</li> </ul>	Azure OpenAI (GPT-4), Autogen prompt chains
<b>4. Data Cleaning</b>	<ul style="list-style-type: none"> <li>• Missing value strategies</li> <li>• Outlier detection &amp; treatment</li> <li>• Type conversions</li> </ul>	pandas, Autogen decision trees
<b>5. Feature Engineering</b>	<ul style="list-style-type: none"> <li>• One-hot / target encoding</li> <li>• Scaling / normalization</li> <li>• Text embeddings (TF-IDF, BERT)</li> <li>• Date/time features</li> </ul>	scikit-learn, Azure ML Feature Store
<b>6. Model Training</b>	<ul style="list-style-type: none"> <li>• Baseline models (linreg, tree)</li> <li>• AutoML runs for quick baselines</li> </ul>	Azure ML, scikit-learn
<b>7. Hyperparameter Tuning</b>	<ul style="list-style-type: none"> <li>• Grid/Bayesian search via HyperDrive</li> <li>• Parallel experiment orchestration</li> </ul>	Azure ML HyperDrive, custom Python
<b>8. Model Evaluation</b>	<ul style="list-style-type: none"> <li>• Metrics (RMSE, AUC, F1)</li> <li>• Confusion matrices &amp; ROC curves</li> <li>• Statistical tests (t-tests, KS-tests)</li> </ul>	scikit-learn, matplotlib for ROC/confusion plots
<b>9. Explainability</b>	<ul style="list-style-type: none"> <li>• SHAP/LIME value computation</li> <li>• Global vs. local explanations</li> <li>• Feature importance plots</li> </ul>	SHAP, LIME, code interpreter + OpenAI for summaries
<b>10. Reporting &amp; Dashboard</b>	<ul style="list-style-type: none"> <li>• Assemble narrative report (Markdown/HTML/PDF)</li> <li>• Embed key static + interactive visuals</li> <li>• Publish to Power BI</li> </ul>	Azure Functions, Power BI REST API, Jinja2
<b>11. Deployment &amp; Monitoring</b>	<ul style="list-style-type: none"> <li>• Containerize model (Docker/ACI)</li> <li>• Expose REST endpoint</li> <li>• Set up telemetry &amp; drift alerts</li> </ul>	Azure Container Instances, App Insights, Monitor
<b>12. Feedback &amp; Iteration</b>	<ul style="list-style-type: none"> <li>• Collect user feedback on dashboards</li> <li>• Trigger retraining or feature tweaks</li> </ul>	Azure Forms, Logic Apps, Autogen loops
<b>13. Knowledge &amp; Memory</b>	<ul style="list-style-type: none"> <li>• Store artifacts &amp; metadata</li> <li>• Answer retrospective queries (“last best model...”)</li> </ul>	Azure Cognitive Search, Cosmos DB

---

### 3. Visualization-Centric Agents

#### 1. Chart Recommendation Agent

- **Input:** DataFrame schema + sample rows
- **Output:** List of suggested chart types (“**histogram**”, “**heatmap**”, “**scatter**”) with rationale

- **How:** Prompt GPT-4 via Autogen with examples of chart-purpose mapping.

## 2. Static Plot Agent

- **Input:** Selected chart spec + data
- **Output:** Matplotlib/Seaborn PNGs
- **How:** Execute code in a sandboxed interpreter, return image URLs.

## 3. Interactive Dashboard Agent

- **Input:** Set of visuals + data endpoints
- **Output:** A live Streamlit/Dash app or embedded Power BI report link
- **How:** Use Azure Functions to deploy the dashboard and publish URL.

## 4. Insight Extraction Agent

- **Input:** Visual outputs + metrics
- **Output:** Natural-language insights (“Correlation between X and Y is strong...”)
- **How:** Use Azure OpenAI to parse chart statistics and draft summaries.

---

# 4. Orchestration Patterns

- **Sequential Pipeline** for full runs (ingest → profile → EDA → ... → deploy).
- **On-Demand Sub-Pipelines:** Analysts can call just “EDA & Visualization” or “Explainability” for new data.
- **Event-Driven:** Feedback Agent triggers a retraining cycle when drift crosses threshold.
- **Parallel Swarms:** Multiple Feature Engineering agents generate candidates concurrently; Orchestrator ranks them.

---

# 5. Next Steps

## 1. Define Agent Contracts

- For each agent, specify JSON schemas: inputs, outputs, tool calls.

## 2. Kick off MVP

- Build Ingestion → Chart Recommendation → Static Plot → Insight Extraction chain.

## 3. Iterate

- Add interactive dashboards and downstream agents one at a time.

## 4. Test & Validate

- Use a sample dataset (e.g., Titanic, NYC taxi) end-to-end.

## 5. Submit

- Fill out your `project.yml`, link code, screenshots of each agent in action, and a short demo GIF.

This design ensures **every** Data Science task—from raw data profiling through advanced visualization, model building, interpretability, deployment, and feedback loops—is handled by a dedicated, testable agent.